

Tying your own hands: credible commitment tools for Frontier Safety Policies

Frontier AI labs make important safety commitments through Frontier Safety Policies (FSP) – but today's FSPs are not robust. Labs have breached or weakened their FSPs, which are structurally flawed in four ways: 1) the consequence of failing is limited; 2) the language is ambiguous; 3) FSPs are easily amended; and 4) few actors exist to uphold FSPs. Labs can and should take three actions to make their FSPs more robust: 1) Set up independent Frontier Safety Oversight Boards; 2) Develop common industry standards via the Frontier Model Forum; and 3) Link financial consequences to safety commitments.

What is the problem?

Frontier AI labs are making important commitments through Frontier Safety Policies (FSPs). FSPs are protocols AI labs write and follow to keep the risks of developing and deploying advanced AI models to an acceptable level. They include critical AI safety commitments around capability thresholds, deployment mitigations, halting plans, and more.

However, many lab commitments have failed. In practice, labs tend not to explicitly breach commitments (though instances exist). Instead, labs often weaken their commitments over time, or simply abandon them.

Existing FSPs are not structurally robust for four reasons. *First, the consequence of failing is limited (P1).* Existing FSP fail to specify consequences for failing commitments, meaning FSPs do not bite and deter undesired behaviour. *Second, the language is ambiguous (P2).* What constitutes a breach can be unclear, and labs can skirt the boundary through subjective judgment. *Third, FSPs are easily amended (P3).* They are framed as “living documents” with loose amendment safeguards, allowing labs to amend away commitments. *Fourth, few actors exist to uphold FSPs (P4).* Powerful pressures and parties push for labs to breach commitment, but few actors have the power or incentive to protect commitments. The underlying driver is competitive pressure across labs, risking *race-to-the-bottom dynamics (D1)*. Existing efforts to strengthen FSPs focus on improving the content of commitments – but these matter little if labs can bypass them structurally.

Why is this important?

FSPs contain substantive commitments. FSPs are arguably the main way through which labs manage AI-enabled harms. Almost all frontier labs have FSPs that they can build upon. Robust FSPs are also in labs' self interest, locking in today's safety preferences against future pressure, and sustaining mutually beneficial co-operation between labs.

FSPs – an example of self-regulation – complement other AI governance channels. Many risk-laden industries combine both self-regulation and government oversight. Over time, voluntary commitments in FSPs can solidify into industry best practices and potentially government policy. Many existing policies and proposals target or rely on FSPs, so robust FSPs support many governance approaches.

But the window for strengthening FSPs is narrowing. OpenAI and Anthropic are planning IPOs, potentially locking in sub-optimal governance structures. New shareholders could worsen incentives to break commitments. As AI capabilities grow, labs could gain greater economic and strategic importance, weakening the ability of existing leverage sources (such as governments, markets, or employees) to install stronger FSPs.

What should we do?

Below are three reinforcing actions that labs can do themselves, today – without the need for new laws, regulations, or government oversight – to strengthen FSPs.

1. Each lab should set up an independent Frontier Safety Oversight Board (FSOB)

Create an actor (P4) to stop labs from exploiting ambiguity (P2) or weakening commitments (P3).

An external FSOB should contract with the lab to carry out two tasks: 1) assess whether the lab is in compliance with its FSP; and 2) oversee the FSP update process. It needs two powers to do so: 1) authoritatively interpret the FSP; and 2) be able to veto changes to the FSP. The FSOB should work closely with the lab itself and third-party auditors to carry out its tasks. The FSOB's mandate, personnel, and funding should be independent from the lab. Violating the FSOB's compliance assessments should come with financial consequences enforced by contracts and markets (see recommendation 3). *An analogue is the Meta Oversight Board.*

2. All labs should develop common industry standards via the Frontier Model Forum (FMF)

Use peer enforcement (P4) to prevent from race-to-the-bottom dynamics on safety (D1)

Labs should task the FMF to develop binding safety standards, or floors, that member lab FSPs must meet. The FMF should certify AI models based on the common standards, providing a soft governance tool other actors can rely on (for example, in procurement or payments). *An analogue is the Institute for Nuclear Power Operations.*

3. Labs should link financial consequences to safety commitments

Supply consequence (P1) and actors (P4) to deter breaches.

Labs should structure financial instruments that impose financial penalties on the lab and align incentives to uphold safety commitments as assessed by the FSOB and the FMF. One way is for labs to issue contingent equity tied to commitments. For example, contingent warrants that dilute existing shareholders, contingent equity grants to lab executives, or contingent golden shares that transfer control. *An analogue is bank regulation.*

Together, these create a basic framework for self-regulation in AI safety. While these actions do not eliminate tail-end risks, they improve on the existing situation, serve as a testing ground and foundation to refine AI governance, and contribute to the defence-in-depth approach to AI-enabled harms. ■